

Computational Intelligence

Topics:

- Learning. Describe the difference between unsupervised and supervised learning including differences in training datasets. Explain the principle of the supervised gradient learning, how the gradient is calculated and hyperparameters of the algorithm.
- Shallow and deep neural networks (DNNs). Typical neurons, their activations and typical activation functions. Describe convolution, pooling and dense layers and their parameters. Application of DNNs for object detection in images.
- Recurrent neural networks (RNN). Explain what the recurrent connections and necessity of their specific evaluation are. Explain fully recurrent and LSTM neural architecture.
- Reinforcement learning. Explain the principle. What is the environmental state, observation, reward, and policy? Describe the actor-critic architecture.**
- Nature-inspired optimization. Explain the principle of the genetic algorithm and the particle swarm optimization algorithm

① Learning: Process by which AI systems improve its performance with experience and data

Supervised	Unsupervised
- Labelled data (P/o pair)	- Unlabelled data
- Predict Outputs for new I/p	- Finds patterns from I/p
- Pys - Reg (continuous o/p)	- Types - Clustering
- Classification [class label]	- grouping similar related data pts
- Algorithm: linear, L1/L2, Polygrad, Regression, KNN, SVM, decision trees.	- Algo: PCA, k-means clustering
- Use: House price forecasting, spam detection	Use: Customer segmentation

Supervised Gradient Learning: Optimizing model's parameters to minimize the cost function and ↑ model's accuracy.

- Optimization by → iteratively adjusting model's params using gradient of a loss (∇f)

Gradient: vector of partial derivatives of loss function w.r.t model parameters $\nabla L(\theta) = \left[\frac{\partial L}{\partial \theta_0}, \frac{\partial L}{\partial \theta_1}, \dots, \frac{\partial L}{\partial \theta_n} \right]$
 points in direction of steepest ascent of function
 In GD, for loss minimization, we move opposite dir of gradient

Gradient Descent Algo: Optimization Algorithm
 hypothesis function $\rightarrow h(\theta) = \theta_0 + \theta_1 x$
 linear reg.

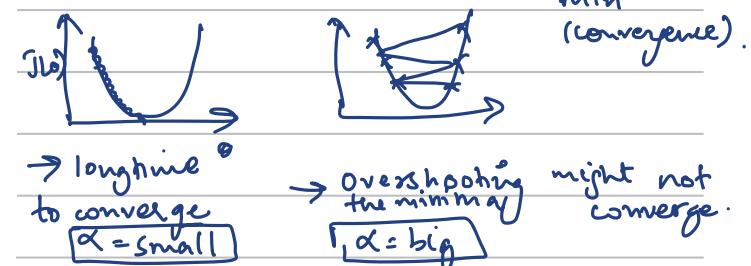
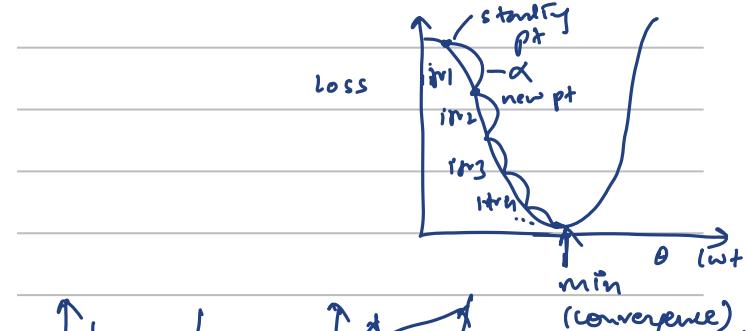
Cost function $\rightarrow J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h(\theta_i) - y_i)^2$
 (MSE fn)
 (LR)
 Goal: minimize $J(\theta_0, \theta_1)$

MSE loss is quadratic

Algorithm:

- Initialize parameters θ
- Calculate gradient of cost() w.r.t parameters

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m (h(\theta_i) - y_i) x_j$$
- update Gradient: $\theta_j := \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j}$
 α = learning rate
 size of step
- Repeat 2 & 3 until cost() converges (i.e. change in $J(\theta)$ becomes very small b/w it) or until preset values of iterations



hyperparameters: α , Epoch, Batchsize

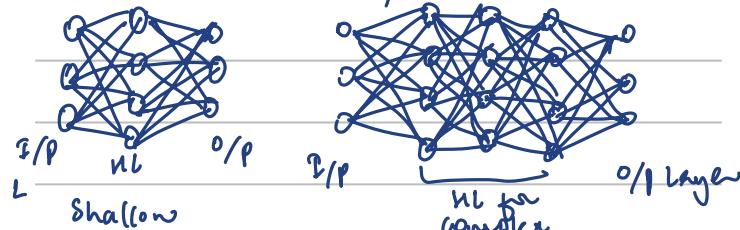
α = determines step size in GD while moving in dir of loss()

Epoch: No. of complete pass through dataset

Batch size: Training examples in 1 iteration

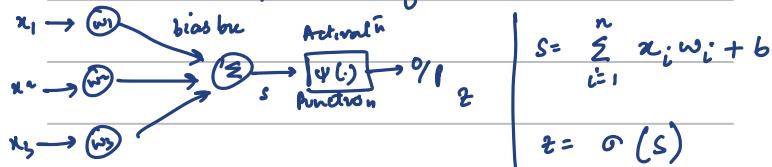
2

Shallow Neural N/w:



- Simple tasks
- Relation b/w I/O is straight forward
- example: Logistic Regression
SL Perceptron
- Complex tasks
- example - Image recognition, NLP, etc
- learns hierarchical features (low-level features - early, high-level (complex) features in late layers)
ex: CNN, RNN

Neuron-Building block of Neural Network



Activation Function: Mathematical function applied to the summation of weight & I/P to introduce non-linearity into the O/P
It helps in learning complex patterns & relationships

Hidden layer: RELU: $f(z) = \begin{cases} 0, & z < 0 \\ z, & z \geq 0 \end{cases}$

AF : CNN
Issue: dying neuron problem because -ve z/p get 0

leaky Relu: $f(z) = \begin{cases} 0.01z, & z < 0 \\ z, & z \geq 0 \end{cases}$

: CNN
: solves dying neuron problem

Tanh : for zero centered activation

RNN

$$\text{tanh} = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

Range (-1, 1)
& mean comes around 0: helpful in hidden layers

Output: Sigmoid: $f(z) = \frac{e^z}{1 + e^z}$

for Binary classification
O/P \rightarrow probabilities

Softmax: Used in Multiclass classification

$$f(z_i) = \frac{e^{z_i}}{\sum e^{z_j}}$$

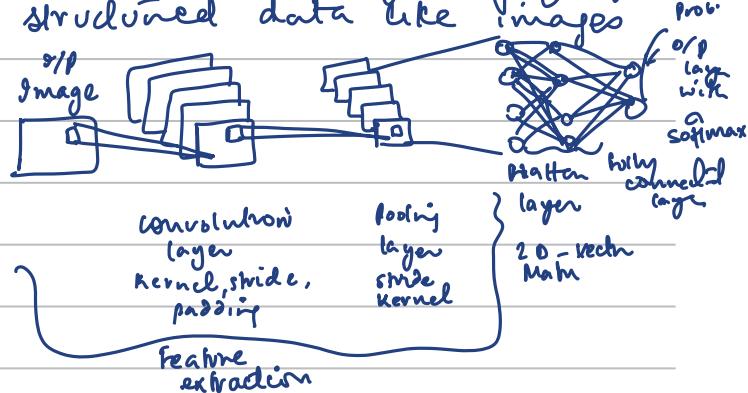
softmax applied to final layer where vector are exponentiated & then normalized by dividing with sum of all e^{z_j}
 \therefore all positive, sum of prob = 1
Range (0, 1)

CNN: Why?

- fewer params
- efficient
- reduced overfitting (\downarrow params)
- complexity Reduced.

Convolutional Neural N/W

\hookrightarrow DL Model specifically for grid structured data like images



S layer: convolutional layer: extract features using kernel/filters to I/P

pooling layer: control dimensions & (max avg) computational load

flatten layer: 2D feature map - 1D feature vector

fully connected: perform further processing

O/P layer: provide final prediction

Kernel: small matrices ($3 \times 3, 5 \times 5$) that slide over the input matrix and perform element wise multiplication & summation to produce feature map.

Stride: No. of pixels by which filter moves across input matrix

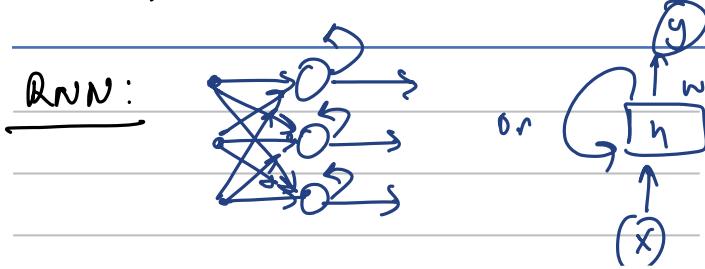
Padding: No. of pixels added to I/P matrix to prevent info loss at the corners along with to control the size of O/P feature map

Output size = $\frac{(\text{Input size} - k + 2P)}{S} + 1$

Pooling: Reduce dimension of fm and preserve important features.

\hookrightarrow Max & Average.

③ RNN, LSTM

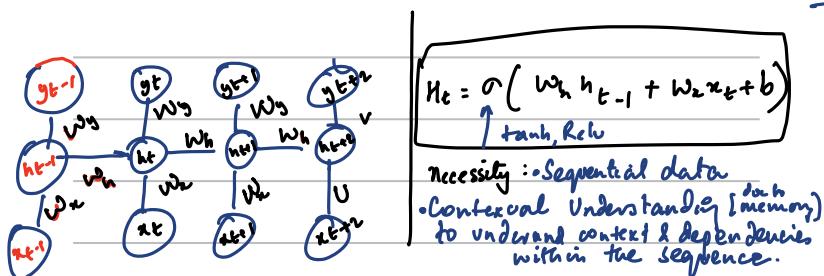


- designed for processing sequential data where current inputs are dependent on the previous steps (NLP, speech, time series)

Hidden state: core

captures the info from previous time stamps. HS is updated at each time stamp as new step is processed allowing N/w to maintain a form of memory ∵ capturing sequential dependencies.

formula for hidden state update:



Issues with RNN: Vanishing gradient (if small data set BP)

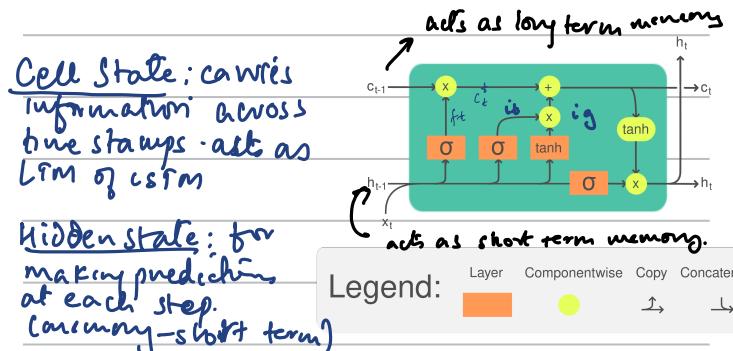
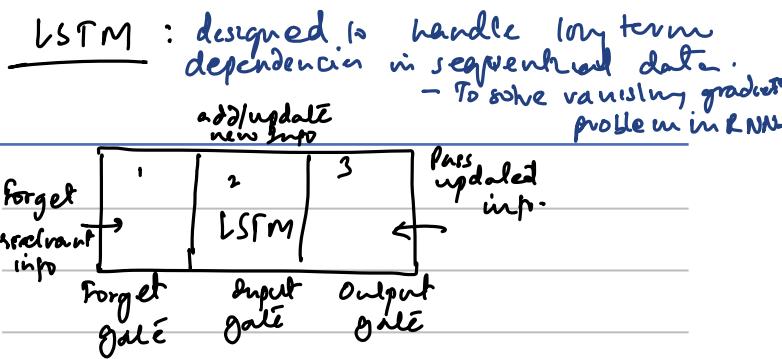
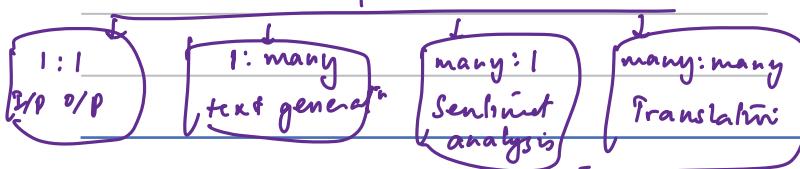
Exploding gradient: (if large)

difficulty with long sequences → LTD.

limited memory (retaining diff over long sequences)

User: NLP, speech recognition, time series prediction, music generation.

RNN types



At each time step,

Forget gate: Use sigmoid (σ) to decide how much previous cell state info to retain

$$f_t = \sigma(w_f x_t + w_{fh} h_{t-1} + b_f)$$

$$c_t = c_{t-1} \odot f_t$$

I/P gate: determines which new info is added to cell state

$$i_t = \sigma(w_i x_t + w_{ih} h_{t-1} + b_i)$$

$$c_t = \tanh(w_c h_{t-1} + w_i x_t + b_c)$$

O/P gate: controls what part of cell state should o/p as hidden state.

σ = determine o/p 0 to 1
 \tanh = scale o/p -1 to 1

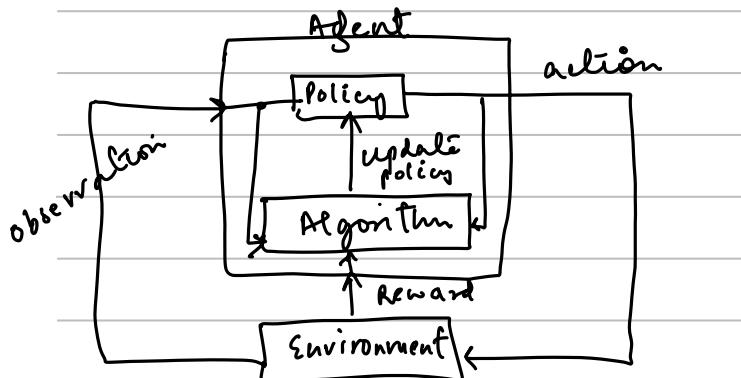
$$o_t = \sigma(w_o x_t + w_{oh} h_{t-1} + b_o)$$

$$h_t = o_t \odot \tanh(c_t)$$

\odot → insutes only relevant part of cell state contributes to o/p

Reinforcement Learning:

Involves agent learning to make decisions by interacting with environment to maximize its cumulative reward w/o human interaction



Agent → decision maker [can]
interacts with environment,
receives observations takes
actions, gets reward.

Environment: world agent interacts in (road, traffic)

Observation: info agent gets from env.
sensor, images camera, GPS.

Action: choice agent makes [turn left,
brake]

Reward: how well agent is doing
+ reward - repeat action
- reward - discourages action

Policy: rule / strategy agent follows
to map observations → actions
It is learned over time.

e.g.: Turn right when board ↑
Stop when red.

Applications of RL:
Self driving cars,
Robots
Recommendation system
Drones

Actor critic Architecture

Type of RL algorithm combines policy + value based methods

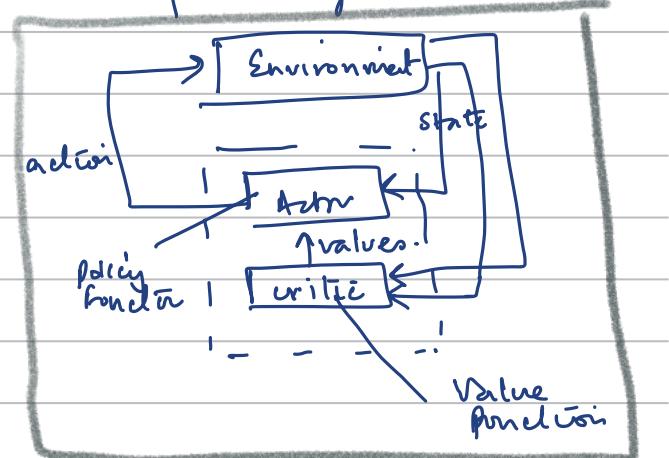
- value based methods: learns value($s \rightarrow V$) that maps each action pair to a value works well when finite actions example → Q learning $\times \rightarrow$ continuous action

policy based methods: for continuous action

directly optimal policy (rule) is found w/o value function.

\times optimal value function diff, how good policy is

Slow learning
reinforce algorithm



Actor: chooses action based on policy
Policy based Method is used
by neural net $\pi(s, a, \theta)$

Critic: Evaluate actions by actor &
and provide fb to improve policy
and - value based method

neural net $\hat{V}(s, a, w)$

When Actor & Critic participate in game,
with time, they start performing better.
learn & become efficient.

Nature Inspired algorithm

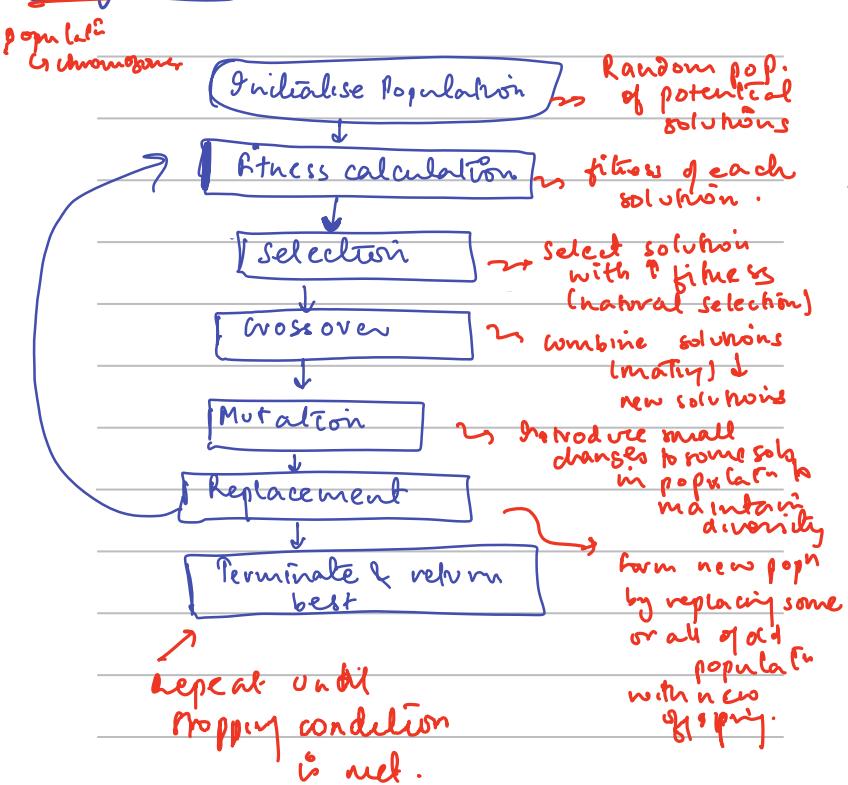
Computational method to solve complex optimisation & search problem by taking inspiration from nature (evolution, natural selection, genetics etc).

Genetic Algorithm:

principle: inspired by Darwin's theory of natural selection & genetics.

Algorithm operates on population of potential solutions, evolving them over generations to find optimum solution to a given problem.

Algorithm:



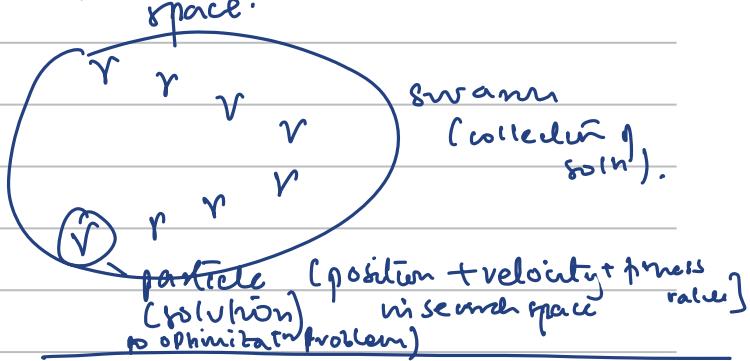
Used to solve : Travel Salesman Problem
& Knapsack Problem

(NP hard problem)
↳ no efficient method exists)

particle Swarm Optimization algorithm.

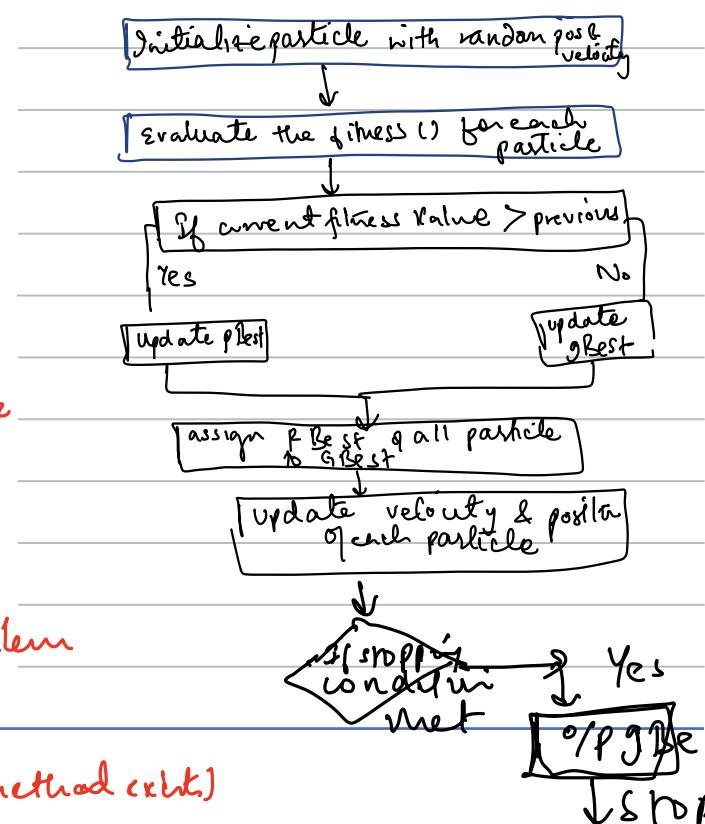
Optimization algo inspired by social behaviour of bird flocking & fish schooling.

Used: to find optimal solution to a problem by simulating swarm of particles moving through search space.



position → potential solution
velocity → direction & speed of movement
fitness value → how good solution is.

particle → keep track of pBest & gBest guiding swarm movement to optimal solution.



Velocity Update formula:

$$v_i(t+1) = \omega v_i(t) + C_1 r_1 (p_{best} - x_i(t)) + C_2 r_2 (g_{best} - x_i(t))$$

personal influence
social influence

$\omega \rightarrow$ inertia weight constant $0 < \omega < 1$
 Diversification controls influence of prev. velocity.

$C_1, C_2 \rightarrow$ cognitive & social coefficient
 control the tradeoff b/w exploration &
 exploitation

Swarm \rightarrow Set of particles (solutions) (S)

Particle \rightarrow a potential solution

Position : $x_i = (x_{i1}, x_{i2}, x_{i3}, \dots, x_{in}) \in \mathbb{R}^n$

Velocity : $v_i = (v_{i1}, v_{i2}, v_{i3}, \dots, v_{in}) \in \mathbb{R}^n$

Each particle maintains : pBest (Individual Best)

Swarm maintains its global best: gBest

Position Update:

$$pos_i(t+1) = pos_i(t) + v_i(t+1)$$