

CAP Theorem

[Brewer's Theorem]

C = Consistency

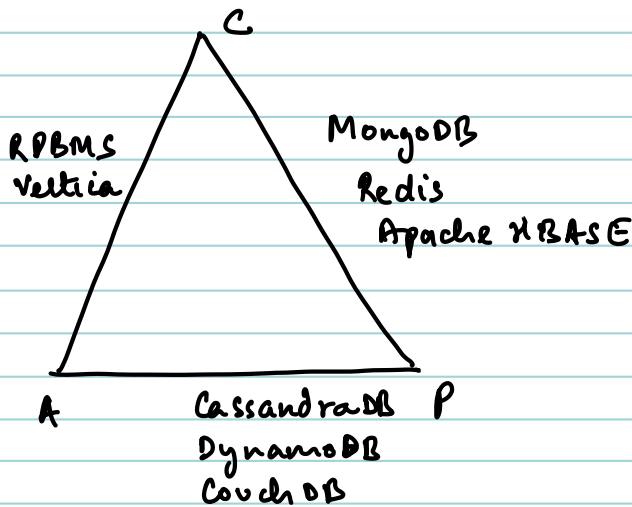
A = Availability

P = Partition Tolerance.



All nodes see same data all the time
→ Each request receives a response (recent or not)

→ System continues to operate despite n/w partitions or connⁿ failures b/w nodes



∴ For a distributed System, It is impossible to achieve all 3 (CAP) simultaneously.

During Partition Tolerance [Network failure], either Consistency or Availability can be provide by a system (DB).

ACID

Atomicity
Consistency
Isolation
Durability

} properly to ensure reliable transaction processing.

Txn is

Atomicity: "All or Nothing"

if any part of txn fails the entire transaction is rolled back.

txn
Account A \rightarrow Account B & no changes are applied.
100 80
 \rightarrow sent 100 80 (Not updated) \rightarrow rolled back

Consistency: follows rules

brings DB from one valid state to another valid state

eg: -ve balance (Rule Violation)
Sum of total remains same

Isolation: Concurrent Txns are isolated from each other.

T_1 | T_2 | T_3
RCA) | RCB) | RCA).

Customer withdrawing money from same account, isolation

|| schedule
to

[Serial
schedule]

ensures that both the txn don't see each other.
 \hookrightarrow It is always consistent.

Durability: All changes in DB should be permanent even if system fails



replication

[SQL]

BASE Properties

NoSQL

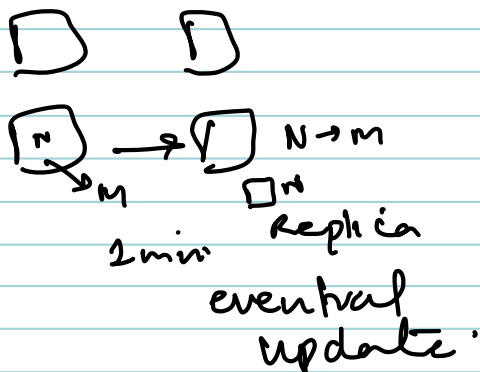
BA → Basically Available
S → Soft State
E → Eventually consistent

Basically Available: System guarantees availability.
DB will respond to all requests even if data sent is stale or incomplete

eg:
Inventory
Item - sold out
not updated
on website
during high
traffic

Soft State

Even w/o Input from outside, the state of system changes.
Due to the eventual updates across distributed system



hw track 1 minute

Eventually
consistent

: System will eventually
become consistent
given enough time.

For updates it will make time (replicas)
But eventually become consistent-