

Jihočeská univerzita
v Českých Budějovicích
University of South Bohemia
in České Budějovice



ústav
aplikované
informatiky



Přírodovědecká
fakulta
Faculty
of Science

ADVANCED DATABASE SYSTEMS

Distributed database systems and NoSQL

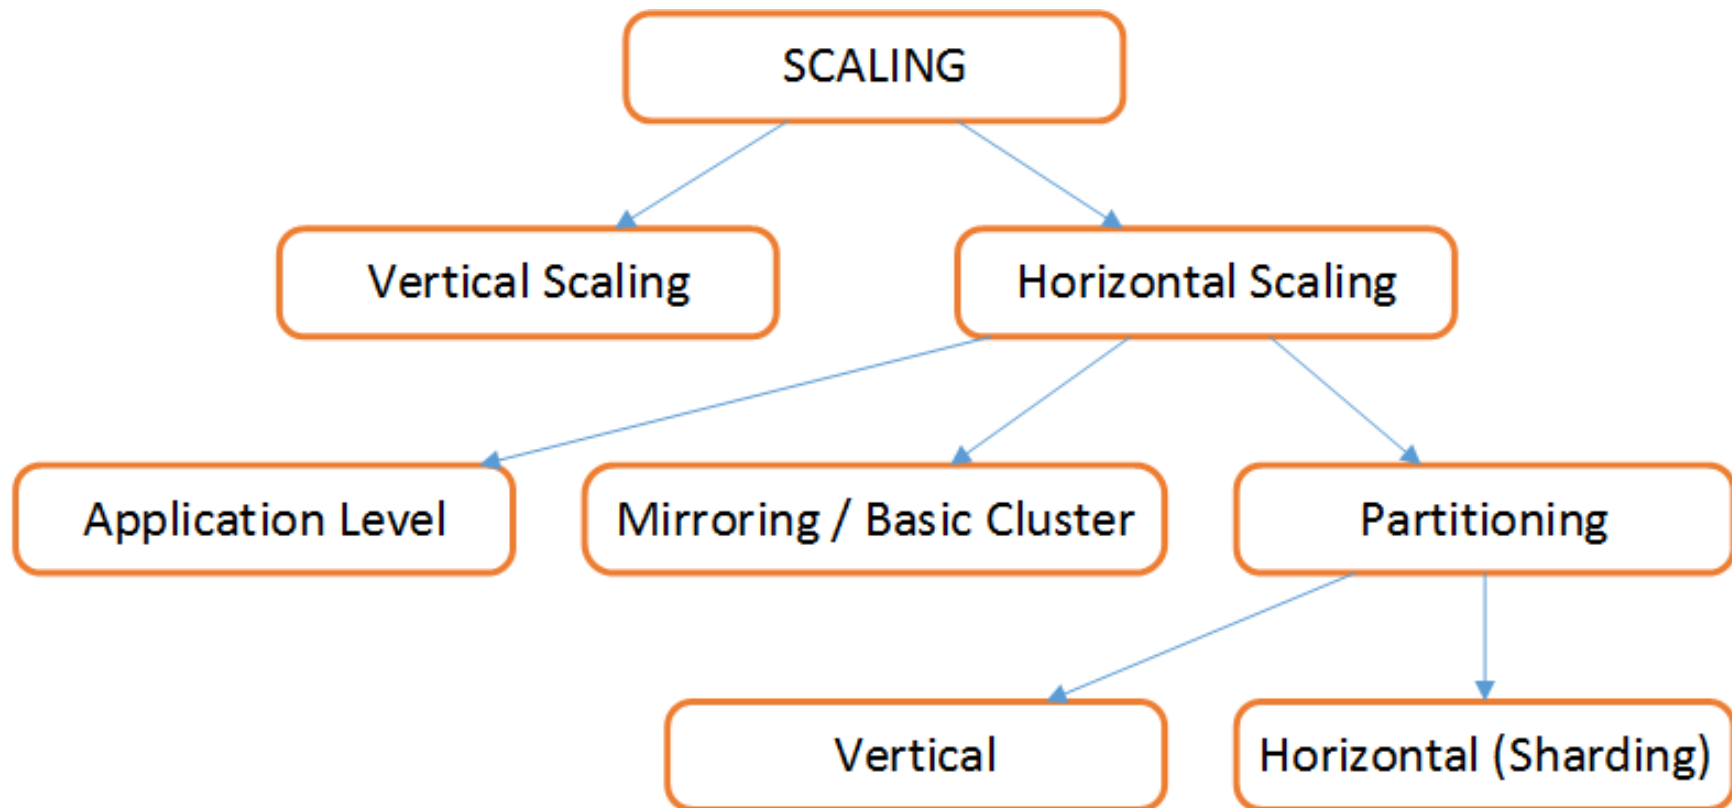
Mgr. Jakub Geyer



➤ Part 2

- Revision
- CAP (Brewer's Theorem)
 - ACID x BASE
 - Transactions
- Practice 3: Indexes
- Practice 4: Document-Oriented database (RavenDB)

Revision



Revision



ústav
aplikované
informatiky



Přírodovědecká
fakulta
Faculty
of Science

- ☐ What are disadvantages/limits of RDBMS?
- ☐ What are the scaling options for databases?
- ☐ What are the horizontal scaling options for database systems?
- ☐ How does the sharding work? What are the common data sharding strategies?
- ☐ What are the characteristics of NoSQL databases?
- ☐ What are 4 types of NoSQL databases? How they work?
Key-Value; Column-Oriented; Document-Oriented; Graph Database

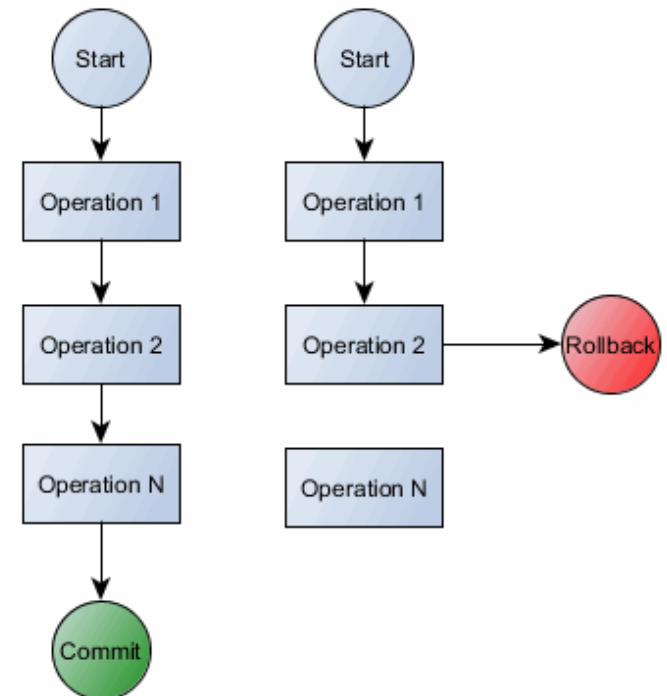
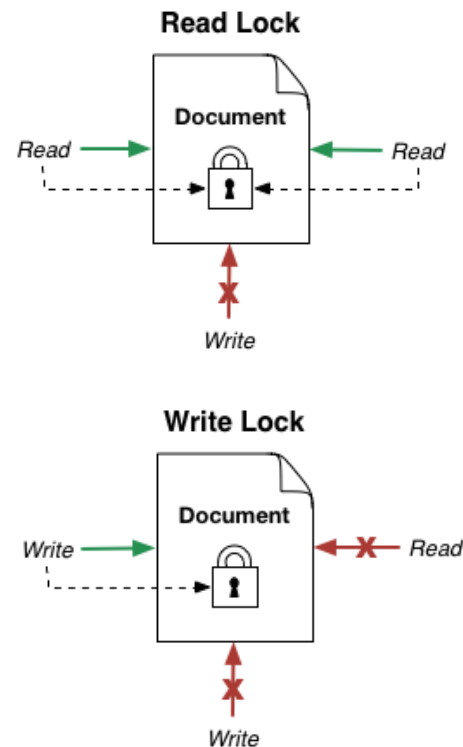
ACID model

Atomicity

Consistency

Isolation

Durability



Which of these features most often do not meet NoSQL databases?

DB transactions



ústav
aplikované
informatiky



Přirodovědecká
fakulta
Faculty
of Science

BASE model

Basic **A**vailability

Soft State

Eventual Consistency

ACID



BASE



It is not true that NoSQL databases cannot meet ACID!

However, be careful about claims that specific NoSQL database meets ACID. They often only meets the requirement for data, not for indexes.

CAP Theorem



ústav
aplikované
informatiky



Přírodovědecká
fakulta
Faculty
of Science

= Brewer's Theorem

CA / CP / AP


Only 2 of 3 are possible at once:

Strong **C**onsistency

High **A**vailability

Partition Tolerance

What belongs here?



This point of view is obsolete!

CAP Theorem




ústav
aplikované
informatiky



Přirodovědecká
fakulta
Faculty
of Science

CAP pro distributed systems:



~~What is the best database? Relational, object-oriented, key-value based, graph?~~

What are the specific requirements?

- Data reliability?
- Working with a large amount of data?
- Read queries speed?
- Lots of writes?
- Complex relationships in data?
- Great data diversity?
- Different views of data? (reporting)
- Transactions?

Indexes



- Significantly affect performance in both RDBMS and NoSQL!!!
- NoSQL databases can often be queried only by indexed values.
- Enables other advanced features (e.g. Map-Reduce, text analyzers, scoring, etc.).

What can usually be indexed in each database type?

(Relational; Key-Value; Column-Oriented; Document-Oriented; Graph Database)





Practice 3:



Advanced indexing and querying using **Apache Lucene**.

<https://lucene.apache.org/core/>

http://lucene.apache.org/core/2_9_4/queryparsersyntax.html

Be careful with *Range* - can work with both numeric values and String (alphabetically).



Practice 4:



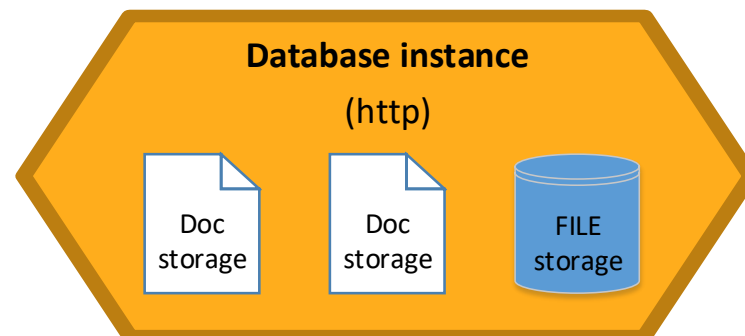
Implementation of simple database using
Document-Oriented database **RavenDB**.

<https://ravendb.net/>

<https://ravendb.net/docs/article-page/3.0/csharp>

<https://azure.microsoft.com/en-us/blog/ravenhq-now-available-in-the-azure-store/>

- Open-Source document-oriented database written in C#.
- Objects (documents) stored as JSON.
- 2 storage types: document storage, file storage
- Storage can be used both separately or as part of a sharding cluster (created and maintained by client application -> no management server required).





Thank you for your attention!